

DATA MINING LAB 2 HOMEWORK REPORT

PRANAV NAGPURE
109065427

Task

Given: A dataset is which was crawled from Twitter, and already labelled with the emotion for these tweets by some specific hashtags in the original text. There are 8 classes (or say emotions) in the dataset: anger, anticipation, disgust, fear, sadness, surprise, trust, and joy.

Objectives: Clean the data by doing some pre-processing first. Then, apply feature engineering or any other data mining technique you have or haven't learned in the Data Mining course. The final goal is to learn a model that can predict the emotion behind each tweet.

1 Pre-processing

1.1 Importing the Data

Tweets in the dataset exist as dictionary with three (key, value) pairs. So, first I create three more columns in the dataframe one for each pair. Then filter the tweets for which emotion is given, i.e., tweets in the training set. However, **for the purpose of evaluating the models I create testing data out of this set later.**

```
data_identification = pd.read_csv('data_identification.csv')
emotion = pd.read_csv('emotion.csv')
tweets = pd.read_json('tweets_DM.json', lines = True)
```


```
tweets['text'] = tweets['_source'].apply(lambda x: x['tweet']['text'])
tweets['hashtags'] = tweets['_source'].apply(lambda x: x['tweet']['hashtags'])
tweets['id'] = tweets['_source'].apply(lambda x: x['tweet']['tweet_id'])
```

1.2 Distribution of classes

I visualized the distribution of classes using a bar graph.

```
tweets_train.emotion.value_counts().plot(kind = 'bar')
```

<AxesSubplot:>



2 Feature Engineering

2.1 Cleaning

We have only one input feature i.e., tweet in text format. And only one output feature i.e., emotion which belong to a finite set of emotions

I perform following operations on the text (in order):

1. Remove all the stop words using neattext library. Stop words are the very common words like 'if', 'but', 'we', 'he', 'she', and 'they'.
2. Remove all the userhandles using neattext library.
3. Remove all punctuations using neattext library.
4. There is '<LH>' in almost every tweet. I remove this using a lambda function.
5. I decapitalize the first character of every word which is not fully capitalized. Fully capitalized words usually suggest that the author of tweet wanted to emphasize on it. So, fully capitalized words remain unchanged. This is done by decap function.
6. Remove all tweets which are blank or became blank after above operations.

Note: The emojis are not affected by these operations and remain unchanged. However, they do not appear in Visualizations. Rather I checked it separately.

```
def decap(text):
    decap_text = ''
    for word in text.split():
        if word.isupper():
            decap_text += word
        else:
            decap_text += word.lower()
    decap_text += ' '
    return decap_text[:-1]
```

```
tweets_train['text'] = tweets_train['text'].apply(nfx.remove_stopwords)
tweets_train['text'] = tweets_train['text'].apply(nfx.remove_userhandles)
tweets_train['text'] = tweets_train['text'].apply(nfx.remove_punctuations)
tweets_train['text'] = tweets_train['text'].apply(lambda x: x.replace('<LH>', ''))
tweets_train['text'] = tweets_train['text'].apply(decap)
```

```
tweets_train = tweets_train[~(tweets_train['text'] == '')]
```

2.2 Converting to Bag Of Words

Converted text to bag of words for all the models (After Visualizations)

```
count_vectorizer = CountVectorizer()
data_X = count_vectorizer.fit_transform(tweets_train['text'])
data_Y = tweets_train['emotion']
X_train, X_test, Y_train, Y_test = train_test_split(data_X, data_Y, test_size = 0.22)
```


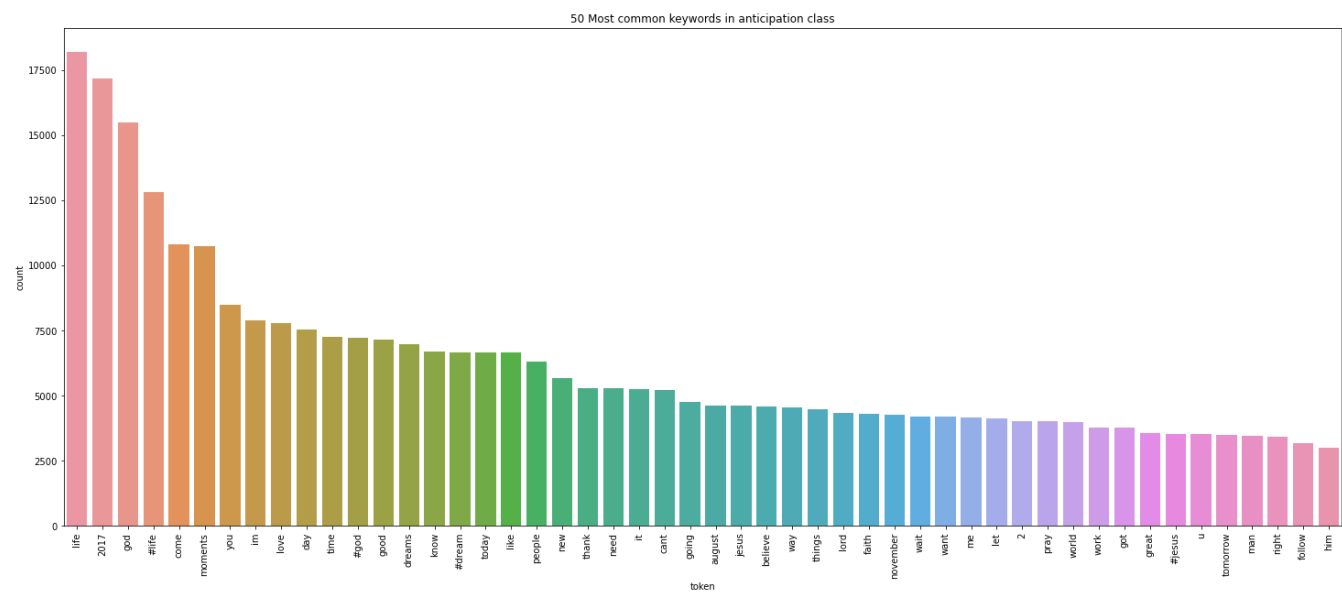
3 Visualizing the Data

3.1 50 Most Common Words by class

Made 8 bar plots one for each emotion (Two examples (Anticipation and Sadness) in next page)

```
def plot_mc_keywords(texts, title):
    keywords = dict(Counter(' '.join(texts.tolist()).split()).most_common(50))
    df = pd.DataFrame(keywords.items(), columns = ['token', 'count'])
    plt.figure(figsize = (25, 10))
    sns.barplot(data = df, x = 'token', y= 'count')
    plt.xticks(rotation = 90)
    plt.title('50 Most common keywords in ' + title + ' class')
    plt.show()

for e in tweets_train.emotion.unique():
    plot_mc_keywords(tweets_train[tweets_train['emotion'] == e]['text'], e)
```



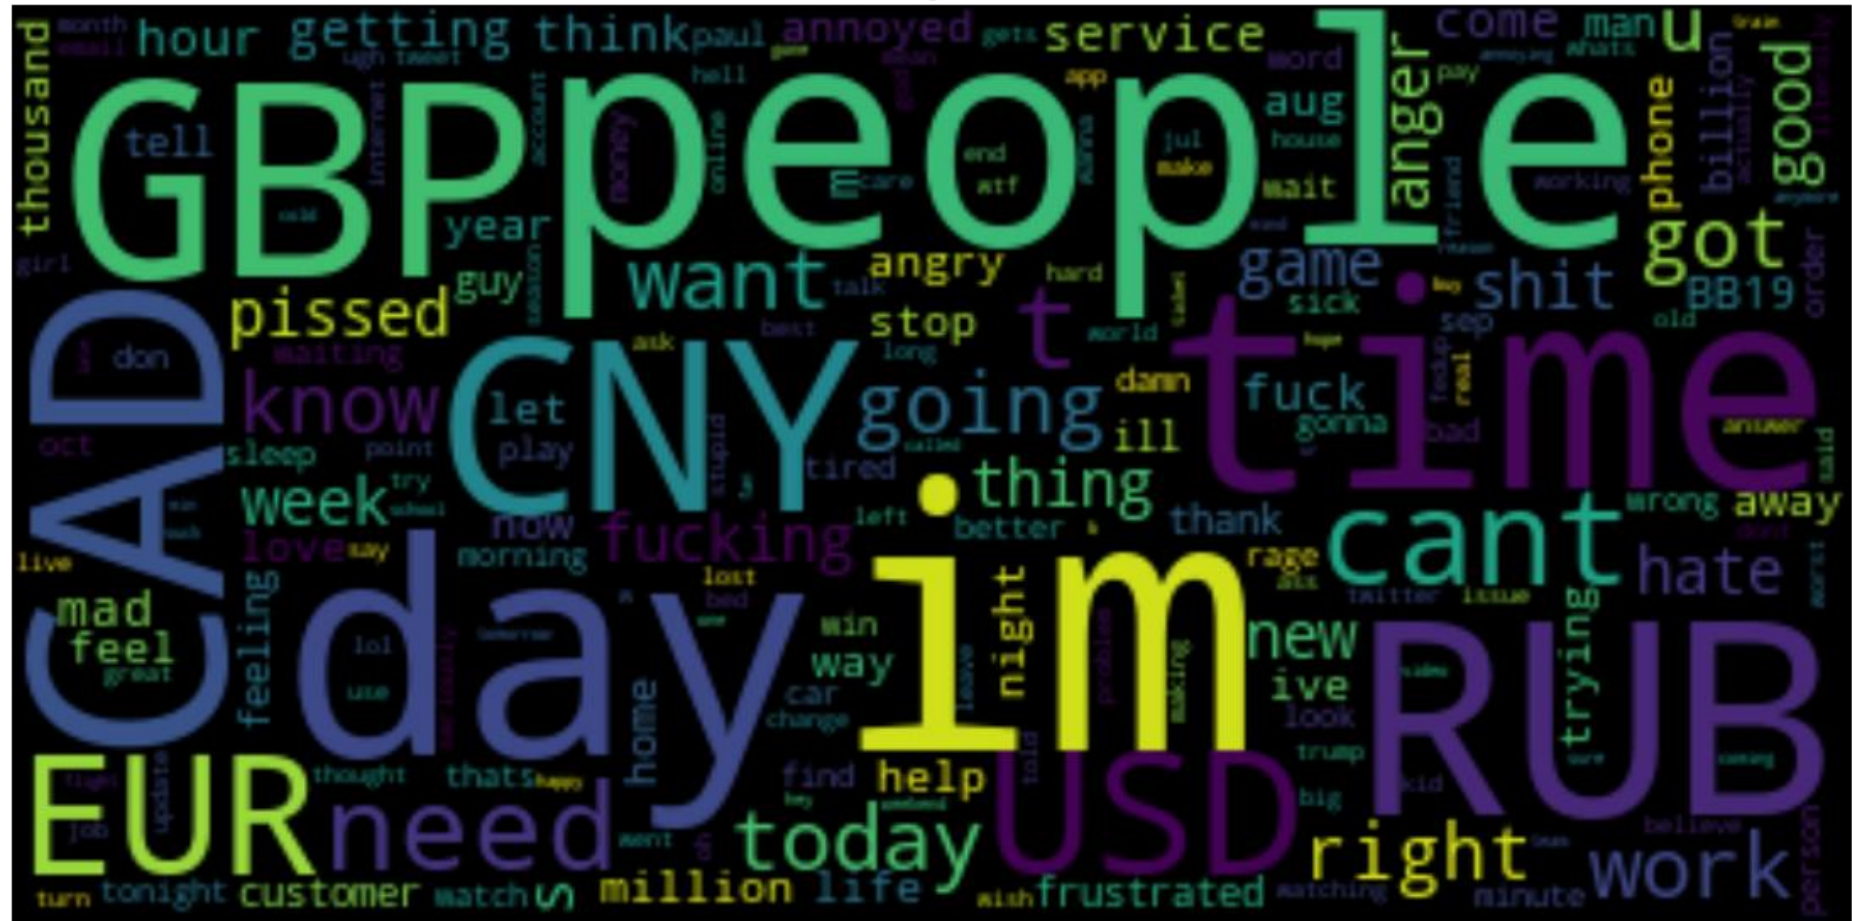
3.2 Word clouds

Made 8 word clouds one for each emotion. These helped to figure out what kind of data cleaning is required. (One example (Anger) in next page)

```
def plot_wordcloud(texts, title):  
    plt.figure(figsize = (25, 10))  
    plt.imshow(WordCloud(collocations = False).generate(' '.join(texts.tolist()))), interpolation = 'bilinear')  
    plt.axis('off')  
    plt.title(title)  
    plt.show()
```

```
for e in tweets_train.emotion.unique():  
    plot_wordcloud(tweets_train[tweets_train['emotion'] == e]['text'], e)
```


anger



4 Naïve Bayes Classifier

Implemented a multinomial naïve bayes classifier, which gives accuracy of 52%. Also, tried complement naïve bayes but it had poor performance compared to multinomial

```
NB_classifier = MultinomialNB()  
NB_classifier.fit(X_train, Y_train)
```

```
MultinomialNB()
```


```
NB_classifier.score(X_test, Y_test)
```

```
0.5203328706648327
```

Evaluation

	precision	recall	f1-score	support
anger	0.81	0.10	0.18	8818
anticipation	0.60	0.52	0.56	54579
disgust	0.47	0.35	0.40	30461
fear	0.74	0.22	0.34	14160
joy	0.51	0.80	0.63	112819
sadness	0.45	0.46	0.45	42539
surprise	0.75	0.12	0.21	10636
trust	0.53	0.26	0.34	45151
accuracy			0.52	319163
macro avg	0.61	0.35	0.39	319163
weighted avg	0.54	0.52	0.49	319163

Very high recall for joy compared to other emotions implies that model is struggling because of skew in distribution



Confucian Matrix for Multinomial Naïve Bayes Classifier

5 Logistic Regression Classifier

Implemented a Logistic Regression classifier, which gives slightly improved accuracy of 54%. Also, tried it with different solvers available. Saga solver gave the best results. It is the recommended solver for high dimensional sparse matrix as per sklearn's documentation.

```
LR_classifier = LogisticRegression(solver = 'saga')
LR_classifier.fit(X_train, Y_train)

C:\Users\prana\anaconda3\lib\site-packages\sklearn\l
means the coef_ did not converge
warnings.warn("The max_iter was reached which mean

LogisticRegression(solver='saga')


LR_classifier.score(X_test, Y_test)

0.5412594818321672
```

Evaluation

	precision	recall	f1-score	support
anger	0.81	0.10	0.18	8818
anticipation	0.60	0.52	0.56	54579
disgust	0.47	0.35	0.40	30461
fear	0.74	0.22	0.34	14160
joy	0.51	0.80	0.63	112819
sadness	0.45	0.46	0.45	42539
surprise	0.75	0.12	0.21	10636
trust	0.53	0.26	0.34	45151
accuracy			0.52	319163
macro avg	0.61	0.35	0.39	319163
weighted avg	0.54	0.52	0.49	319163

Once again, very high recall for joy compared to other emotions implies that model is struggling because of skew in distribution



Confucian Matrix for Multinomial Naïve Bayes Classifier